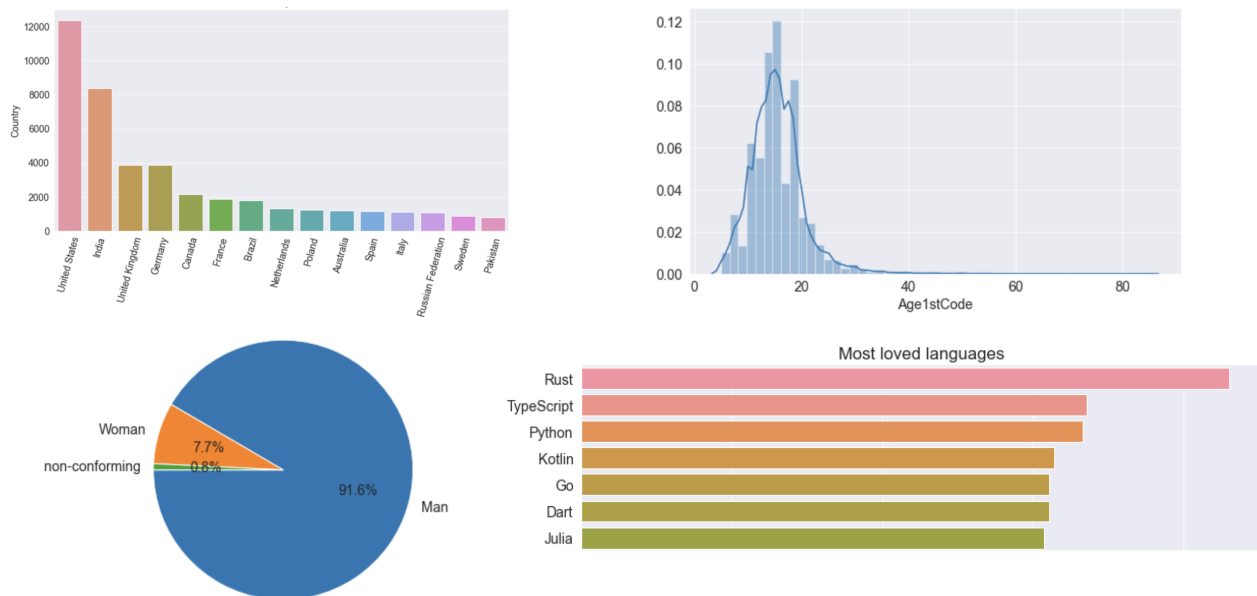


Exploratory Data Analysis using Python - A Case Study

Analyzing responses from the Stack Overflow Annual Developer Survey 2020



Part 9 of "Data Analysis with Python: Zero to Pandas"

This tutorial series is a beginner-friendly introduction to programming and data analysis using the Python programming language. These tutorials take a practical and coding-focused approach. The best way to learn the material is to execute the code and experiment with it yourself. Check out the full series here:

1. [First Steps with Python and Jupyter](#)
2. [A Quick Tour of Variables and Data Types](#)
3. [Branching using Conditional Statements and Loops](#)
4. [Writing Reusable Code Using Functions](#)
5. [Reading from and Writing to Files](#)
6. [Numerical Computing with Python and Numpy](#)
7. [Analyzing Tabular Data using Pandas](#)
8. [Data Visualization using Matplotlib & Seaborn](#)
9. [Exploratory Data Analysis - A Case Study](#)

The following topics are covered in this tutorial:

- Selecting and downloading a dataset
- Data preparation and cleaning
- Exploratory analysis and visualization
- Asking and answering interesting questions
- Summarizing inferences and drawing conclusions

How to run the code

This tutorial is an executable [Jupyter notebook](#) hosted on [Jovian](#). You can *run* this tutorial and experiment with the code examples in a couple of ways: *using free online resources* (recommended) or *on your computer*.

Option 1: Running using free online resources (1-click, recommended)

The easiest way to start executing the code is to click the **Run** button at the top of this page and select **Run on Binder**. You can also select "Run on Colab" or "Run on Kaggle", but you'll need to create an account on [Google Colab](#) or [Kaggle](#) to use these platforms.

Option 2: Running on your computer locally

To run the code on your computer locally, you'll need to set up [Python](#), download the notebook and install the required libraries. We recommend using the [Conda](#) distribution of Python. Click the **Run** button at the top of this page, select the **Run Locally** option, and follow the instructions.

Jupyter Notebooks: This tutorial is a [Jupyter notebook](#) - a document made of *cells*. Each cell can contain code written in Python or explanations in plain English. You can execute code cells and view the results, e.g., numbers, messages, graphs, tables, files, etc., instantly within the notebook. Jupyter is a powerful platform for experimentation and analysis. Don't be afraid to mess around with the code & break things - you'll learn a lot by encountering and fixing errors. You can use the "Kernel > Restart & Clear Output" menu option to clear all outputs and start again from the top.

Introduction

In this tutorial, we'll analyze the StackOverflow developer survey dataset. The dataset contains responses to an annual survey conducted by StackOverflow. You can find the raw data & official analysis here:

<https://insights.stackoverflow.com/survey>.

There are several options for getting the dataset into Jupyter:

- Download the CSV manually and upload it via Jupyter's GUI
- Use the `urlretrieve` function from the `urllib.request` to download CSV files from a raw URL
- Use a helper library, e.g., [opendatasets](#), which contains a collection of curated datasets and provides a helper function for direct download.

We'll use the `opendatasets` helper library to download the files.

```
import opendatasets as od
```

```
od.download('stackoverflow-developer-survey-2020')
```

Downloading

https://raw.githubusercontent.com/JovianML/opendatasets/master/data/stackoverflow-developer-survey-2020/survey_results_public.csv to `./stackoverflow-developer-survey-2020/survey_results_public.csv`

100%|██████████| 94576640/94603888 [06:48<00:00, 126126.63it/s]

0it [00:00, ?it/s]

Downloading

https://raw.githubusercontent.com/JovianML/opendatasets/master/data/stackoverflow-developer-survey-2020/survey_results_schema.csv to ./stackoverflow-developer-survey-2020/survey_results_schema.csv

0%| | 0/8428 [00:01<?, ?it/s]

0it [00:00, ?it/s]

Downloading

<https://raw.githubusercontent.com/JovianML/opendatasets/master/data/stackoverflow-developer-survey-2020/README.txt> to ./stackoverflow-developer-survey-2020/README.txt

0%| | 0/2268 [00:00<?, ?it/s]

Let's verify that the dataset was downloaded into the directory `stackoverflow-developer-survey-2020` and retrieve the list of files in the dataset.

```
import os
```

```
os.listdir('stackoverflow-developer-survey-2020')
```

```
['survey_results_public.csv', 'survey_results_schema.csv', 'README.txt']
```

94609408it [07:00, 126126.63it/s]

You can through the downloaded files using the "File" > "Open" menu option in Jupyter. It seems like the dataset contains three files:

- `README.txt` - Information about the dataset
- `survey_results_schema.csv` - The list of questions, and shortcodes for each question
- `survey_results_public.csv` - The full list of responses to the questions

Let's load the CSV files using the Pandas library. We'll use the name `survey_raw_df` for the data frame to indicate this is unprocessed data that we might clean, filter, and modify to prepare a data frame ready for analysis.

```
import pandas as pd
```

16384it [00:25, 630.24it/s]

8192it [00:24, 330.41it/s]

94609408it [07:14, 217588.78it/s]

```
survey_raw_df = pd.read_csv('stackoverflow-developer-survey-2020/survey_results_public.
```

```
survey_raw_df
```

	Respondent	MainBranch	Hobbyist	Age	Age1stCode	CompFreq	CompTotal	ConvertedComp	Country	C
0	1	I am a developer by profession	Yes	NaN	13	Monthly	NaN	NaN	Germany	
1	2	I am a developer by profession	No	NaN	19	NaN	NaN	NaN	United Kingdom	
2	3	I code primarily as a hobby	Yes	NaN	15	NaN	NaN	NaN	Russian Federation	
3	4	I am a developer by profession	Yes	25.0	18	NaN	NaN	NaN	Albania	
4	5	I used to be a developer by profession, but no...	Yes	31.0	16	NaN	NaN	NaN	United States	
...	
64456	64858	NaN	Yes	NaN	16	NaN	NaN	NaN	United States	
64457	64867	NaN	Yes	NaN	NaN	NaN	NaN	NaN	Morocco	
64458	64898	NaN	Yes	NaN	NaN	NaN	NaN	NaN	Viet Nam	
64459	64925	NaN	Yes	NaN	NaN	NaN	NaN	NaN	Poland	
64460	65112	NaN	Yes	NaN	NaN	NaN	NaN	NaN	Spain	

64461 rows × 61 columns

The dataset contains over 64,000 responses to 60 questions (although many questions are optional). The responses have been anonymized to remove personally identifiable information, and each respondent has been assigned a randomized respondent ID.

Let's view the list of columns in the data frame.

```
survey_raw_df.columns
```

```
Index(['Respondent', 'MainBranch', 'Hobbyist', 'Age', 'Age1stCode', 'CompFreq',  
      'CompTotal', 'ConvertedComp', 'Country', 'CurrencyDesc',  
      'CurrencySymbol', 'DatabaseDesireNextYear', 'DatabaseWorkedWith',  
      'DevType', 'EdLevel', 'Employment', 'Ethnicity', 'Gender', 'JobFactors',
```

```
'JobSat', 'JobSeek', 'LanguageDesireNextYear', 'LanguageWorkedWith',
'MiscTechDesireNextYear', 'MiscTechWorkedWith',
'NEWCollabToolsDesireNextYear', 'NEWCollabToolsWorkedWith', 'NEWDevOps',
'NEWDevOpsImpt', 'NEWEdImpt', 'NEWJobHunt', 'NEWJobHuntResearch',
'NEWLearn', 'NEWOffTopic', 'NEWOnboardGood', 'NEWOtherComms',
'NEWOvertime', 'NEWPurchaseResearch', 'NEWPurpleLink', 'NEWSOSites',
'NEWStuck', 'OpSys', 'OrgSize', 'PlatformDesireNextYear',
'PlatformWorkedWith', 'PurchaseWhat', 'Sexuality', 'SOAccount',
'SOComm', 'SOPartFreq', 'SOVisitFreq', 'SurveyEase', 'SurveyLength',
'Trans', 'UndergradMajor', 'WebframeDesireNextYear',
'WebframeWorkedWith', 'WelcomeChange', 'WorkWeekHrs', 'YearsCode',
'YearsCodePro'],
dtype='object')
```

It appears that shortcodes for questions have been used as column names.

We can refer to the schema file to see the full text of each question. The schema file contains only two columns: `Column` and `QuestionText`. We can load it as Pandas Series with `Column` as the index and the `QuestionText` as the value.

```
schema_fname = 'stackoverflow-developer-survey-2020/survey_results_schema.csv'
schema_raw = pd.read_csv(schema_fname, index_col='Column').QuestionText
```

schema_raw

```
Column
Respondent      Randomized respondent ID number (not in order ...
MainBranch      Which of the following options best describes ...
Hobbyist        Do you code as a hobby?
Age             What is your age (in years)? If you prefer not...
Age1stCode      At what age did you write your first line of c...
...
WebframeWorkedWith  Which web frameworks have you done extensive d...
WelcomeChange     Compared to last year, how welcome do you feel...
WorkWeekHrs       On average, how many hours per week do you wor...
YearsCode        Including any education, how many years have y...
YearsCodePro      NOT including education, how many years have y...
Name: QuestionText, Length: 61, dtype: object
```

We can now use `schema_raw` to retrieve the full question text for any column in `survey_raw_df`.

```
schema_raw['YearsCodePro']
```

```
'NOT including education, how many years have you coded professionally (as a part of
your work)?'
```

We've now loaded the dataset. We're ready to move on to the next step of preprocessing & cleaning the data for our analysis.

Save and upload your notebook

Whether you're running this Jupyter notebook online or on your computer, it's essential to save your work from time to time. You can continue working on a saved notebook later or share it with friends and colleagues to let them execute your code. [Jovian](#) offers an easy way of saving and sharing your Jupyter notebooks online.

```
# Select a project name
project='python-eda-stackoverflow-survey'
```

```
# Install the Jovian library
!pip install jovian --upgrade --quiet
```

```
import jovian
```

```
jovian.commit(project=project)
```

```
[jovian] Attempting to save notebook..
[jovian] Updating notebook "aakashns/python-eda-stackoverflow-survey" on
https://jovian.ai/
[jovian] Uploading notebook..
[jovian] Capturing environment..
[jovian] Committed successfully! https://jovian.ai/aakashns/python-eda-stackoverflow-survey
'https://jovian.ai/aakashns/python-eda-stackoverflow-survey'
```

The first time you run `jovian.commit`, you'll be asked to provide an API Key to securely upload the notebook to your Jovian account. You can get the API key from your [Jovian profile page](#) after logging in / signing up.

`jovian.commit` uploads the notebook to your Jovian account, captures the Python environment, and creates a shareable link for your notebook, as shown above. You can use this link to share your work and let anyone (including you) run your notebooks and reproduce your work.

Data Preparation & Cleaning

While the survey responses contain a wealth of information, we'll limit our analysis to the following areas:

- Demographics of the survey respondents and the global programming community
- Distribution of programming skills, experience, and preferences
- Employment-related information, preferences, and opinions

Let's select a subset of columns with the relevant data for our analysis.

```
selected_columns = [
    # Demographics
    'Country',
    'Age',
    'Gender',
    'EdLevel',
    'UndergradMajor',
```

```

# Programming experience
'Hobbyist',
'Age1stCode',
'YearsCode',
'YearsCodePro',
'LanguageWorkedWith',
'LanguageDesireNextYear',
'NEWLearn',
'NEWStuck',
# Employment
'Employment',
'DevType',
'WorkWeekHrs',
'JobSat',
'JobFactors',
'NEWOvertime',
'NEWEdImpt'
]

```

```
len(selected_columns)
```

20

Let's extract a copy of the data from these columns into a new data frame `survey_df` . We can continue to modify further without affecting the original data frame.

```
survey_df = survey_raw_df[selected_columns].copy()
```

```
schema = schema_raw[selected_columns]
```

Let's view some basic information about the data frame.

```
survey_df.shape
```

(64461, 20)

```
survey_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 64461 entries, 0 to 64460
```

```
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	Country	64072 non-null	object
1	Age	45446 non-null	float64
2	Gender	50557 non-null	object
3	EdLevel	57431 non-null	object
4	UndergradMajor	50995 non-null	object

5	Hobbyist	64416	non-null	object
6	Age1stCode	57900	non-null	object
7	YearsCode	57684	non-null	object
8	YearsCodePro	46349	non-null	object
9	LanguageWorkedWith	57378	non-null	object
10	LanguageDesireNextYear	54113	non-null	object
11	NEWLearn	56156	non-null	object
12	NEWStuck	54983	non-null	object
13	Employment	63854	non-null	object
14	DevType	49370	non-null	object
15	WorkWeekHrs	41151	non-null	float64
16	JobSat	45194	non-null	object
17	JobFactors	49349	non-null	object
18	NEWOvertime	43231	non-null	object
19	NEWEdImpt	48465	non-null	object

```
dtypes: float64(2), object(18)
memory usage: 9.8+ MB
```

Most columns have the data type `object` , either because they contain values of different types or contain empty values (`NaN`). It appears that every column contains some empty values since the Non-Null count for every column is lower than the total number of rows (64461). We'll need to deal with empty values and manually adjust the data type for each column on a case-by-case basis.

Only two of the columns were detected as numeric columns (`Age` and `WorkWeekHrs`), even though a few other columns have mostly numeric values. To make our analysis easier, let's convert some other columns into numeric data types while ignoring any non-numeric value. The non-numeric are converted to `NaN` .

```
survey_df['Age1stCode'] = pd.to_numeric(survey_df.Age1stCode, errors='coerce')
survey_df['YearsCode'] = pd.to_numeric(survey_df.YearsCode, errors='coerce')
survey_df['YearsCodePro'] = pd.to_numeric(survey_df.YearsCodePro, errors='coerce')
```

Let's now view some basic statistics about numeric columns.

```
survey_df.describe()
```

	Age	Age1stCode	YearsCode	YearsCodePro	WorkWeekHrs
count	45446.000000	57473.000000	56784.000000	44133.000000	41151.000000
mean	30.834111	15.476572	12.782051	8.869667	40.782174
std	9.585392	5.114081	9.490657	7.759961	17.816383
min	1.000000	5.000000	1.000000	1.000000	1.000000
25%	24.000000	12.000000	6.000000	3.000000	40.000000
50%	29.000000	15.000000	10.000000	6.000000	40.000000
75%	35.000000	18.000000	17.000000	12.000000	44.000000
max	279.000000	85.000000	50.000000	50.000000	475.000000

There seems to be a problem with the age column, as the minimum value is 1 and the maximum is 279. This is a common issue with surveys: responses may contain invalid values due to accidental or intentional errors while responding. A simple fix would be to ignore the rows where the age is higher than 100 years or lower than 10 years as invalid survey responses. We can do this using the `.drop` method, [as explained here](#).

```
survey_df.drop(survey_df[survey_df.Age < 10].index, inplace=True)
survey_df.drop(survey_df[survey_df.Age > 100].index, inplace=True)
```

The same holds for `WorkWeekHrs`. Let's ignore entries where the value for the column is higher than 140 hours. (~20 hours per day).

```
survey_df.drop(survey_df[survey_df.WorkWeekHrs > 140].index, inplace=True)
```

The gender column also allows for picking multiple options. We'll remove values containing more than one option to simplify our analysis.

```
survey_df['Gender'].value_counts()
```

```
Man                                45895
Woman                             3835
Non-binary, genderqueer, or gender non-conforming    385
Man;Non-binary, genderqueer, or gender non-conforming    121
Woman;Non-binary, genderqueer, or gender non-conforming    92
Woman;Man                                73
Woman;Man;Non-binary, genderqueer, or gender non-conforming    25
Name: Gender, dtype: int64
```

```
import numpy as np
```

```
survey_df.where(~(survey_df.Gender.str.contains('; ', na=False)), np.nan, inplace=True)
```

We've now cleaned up and prepared the dataset for analysis. Let's take a look at a sample of rows from the data frame.

```
survey_df.sample(10)
```

	Country	Age	Gender	EdLevel	UndergradMajor	Hobbyist	Age1stCode	YearsCode	YearsCodePro
29466	Germany	25.0	Man	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	Another engineering discipline (such as civil,...	Yes	15.0	4.0	NaN
39432	United States	NaN	NaN	NaN	NaN	Yes	NaN	NaN	NaN
45451	Belgium	29.0	Man	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Web development or web design	Yes	18.0	10.0	6.0

	Country	Age	Gender	EdLevel	UndergradMajor	Hobbyist	Age1stCode	YearsCode	YearsCodePro
63318	Viet Nam	NaN	NaN	Associate degree (A.A., A.S., etc.)	A business discipline (such as accounting, fin...	Yes	27.0	2.0	2.0
29520	United States	33.0	Man	Other doctoral degree (Ph.D., Ed.D., etc.)	Another engineering discipline (such as civil,...	Yes	15.0	15.0	10.0
38253	Czech Republic	NaN	NaN	Some college/university study without earning ...	I never declared a major	Yes	12.0	26.0	9.0
58351	India	NaN	NaN	NaN	NaN	Yes	NaN	NaN	NaN
52082	United States	33.0	Man	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Computer science, computer engineering, or sof...	Yes	13.0	20.0	10.0
47928	Sweden	30.0	Man	Bachelor's degree (B.A., B.S., B.Eng., etc.)	Computer science, computer engineering, or sof...	Yes	15.0	9.0	5.0
2340	United States	36.0	Man	Some college/university study without earning ...	Another engineering discipline (such as civil,...	No	30.0	36.0	4.0

Let's save and commit our work before continuing.

```
import jovian
```

```
jovian.commit()
```

```
[jovian] Attempting to save notebook..
```

```
[jovian] Updating notebook "aakashns/python-eda-stackoverflow-survey" on
```

```
https://jovian.ai/
```

```
[jovian] Uploading notebook..
```

```
[jovian] Capturing environment..
```

```
[jovian] Committed successfully! https://jovian.ai/aakashns/python-eda-stackoverflow-survey
```

```
'https://jovian.ai/aakashns/python-eda-stackoverflow-survey'
```

Exploratory Analysis and Visualization

Before we ask questions about the survey responses, it would help to understand the respondents' demographics, i.e., country, age, gender, education level, employment level, etc. It's essential to explore these variables to understand how representative the survey is of the worldwide programming community. A survey of this scale generally tends to have some [selection bias](#).

Let's begin by importing `matplotlib.pyplot` and `seaborn`.

```
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

sns.set_style('darkgrid')
matplotlib.rcParams['font.size'] = 14
matplotlib.rcParams['figure.figsize'] = (9, 5)
matplotlib.rcParams['figure.facecolor'] = '#00000000'
```

Country

Let's look at the number of countries from which there are responses in the survey and plot the ten countries with the highest number of responses.

```
schema.Country
```

```
'Where do you live?'
```

```
survey_df.Country.nunique()
```

```
183
```

We can identify the countries with the highest number of respondents using the `value_counts` method.

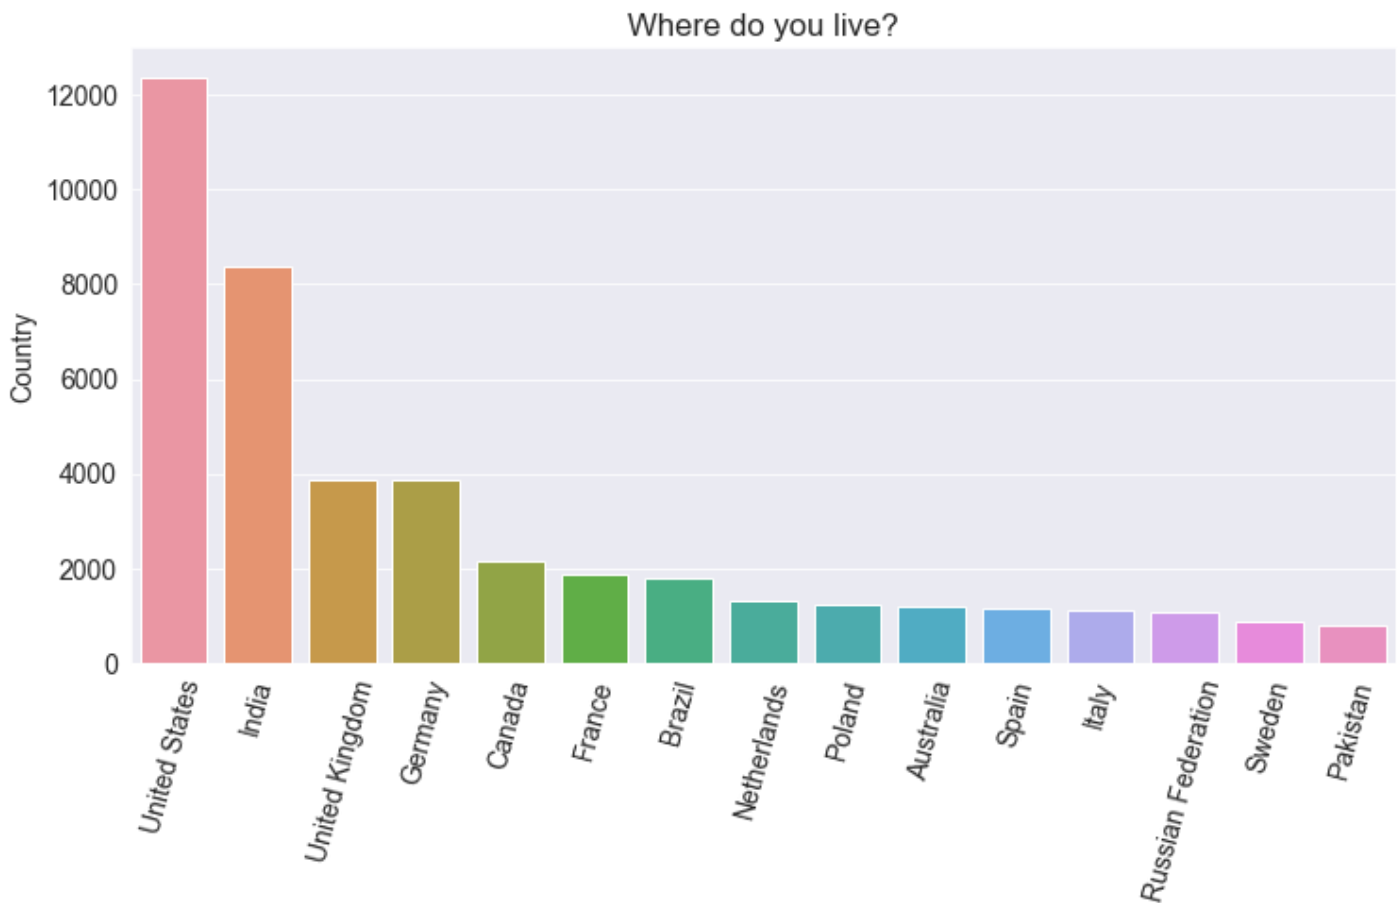
```
top_countries = survey_df.Country.value_counts().head(15)
top_countries
```

United States	12371
India	8364
United Kingdom	3881
Germany	3864
Canada	2175
France	1884
Brazil	1804
Netherlands	1332
Poland	1259
Australia	1199
Spain	1157
Italy	1115
Russian Federation	1085
Sweden	879
Pakistan	802

Name: Country, dtype: int64

We can visualize this information using a bar chart.

```
plt.figure(figsize=(12,6))
plt.xticks(rotation=75)
plt.title(schema.Country)
sns.barplot(x=top_countries.index, y=top_countries);
```



It appears that a disproportionately high number of respondents are from the US and India, probably because the survey is in English, and these countries have the highest English-speaking populations. We can already see that the survey may not be representative of the global programming community - especially from non-English speaking countries. Programmers from non-English speaking countries are almost certainly underrepresented.

Exercise: Try finding the percentage of responses from English-speaking vs. non-English speaking countries. You can use [this list of languages spoken in different countries](#).

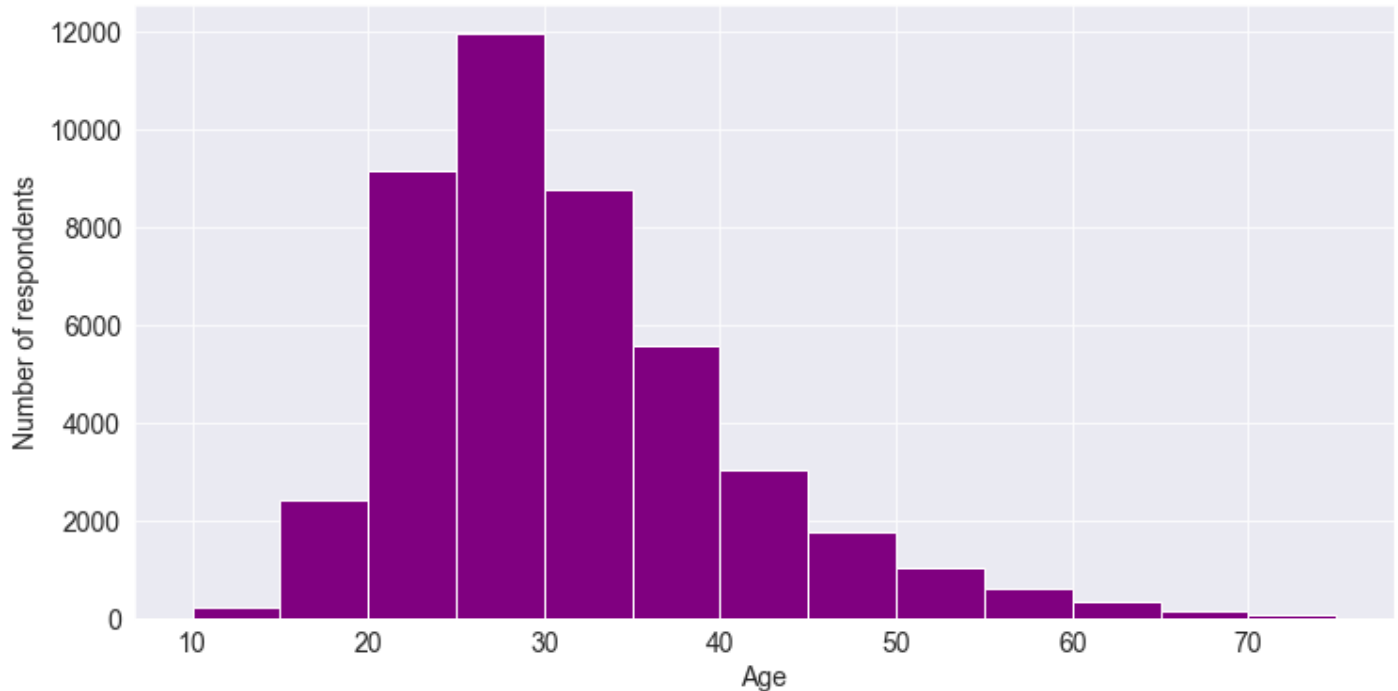
Age

The distribution of respondents' age is another crucial factor to look at. We can use a histogram to visualize it.

```
plt.figure(figsize=(12, 6))
plt.title(schema.Age)
plt.xlabel('Age')
plt.ylabel('Number of respondents')

plt.hist(survey_df.Age, bins=np.arange(10,80,5), color='purple');
```

What is your age (in years)? If you prefer not to answer, you may leave this question blank.



It appears that a large percentage of respondents are 20-45 years old. It's somewhat representative of the programming community in general. Many young people have taken up computer science as their field of study or profession in the last 20 years.

Exercise: You may want to filter out responses by age (or age group) if you'd like to analyze and compare the survey results for different age groups. Create a new column called `AgeGroup` containing values like `Less than 10 years`, `10-18 years`, `18-30 years`, `30-45 years`, `45-60 years` and `Older than 60 years`. Then, repeat the analysis in the rest of this notebook for each age group.

Gender

Let's look at the distribution of responses for the Gender. It's a well-known fact that women and non-binary genders are underrepresented in the programming community, so we might expect to see a skewed distribution here.

```
schema.Gender
```

'Which of the following describe you, if any? Please check all that apply. If you prefer not to answer, you may leave this question blank.'

```
gender_counts = survey_df.Gender.value_counts()
gender_counts
```

Man	45895
Woman	3835
Non-binary, genderqueer, or gender non-conforming	385

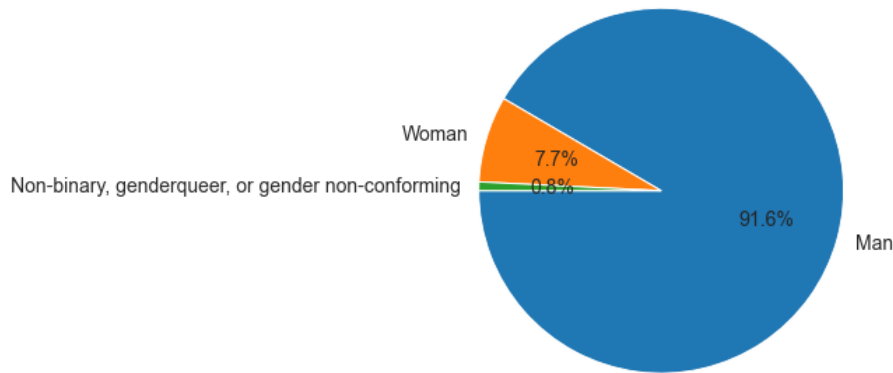
Name: Gender, dtype: int64

A pie chart would be a great way to visualize the distribution.

```
plt.figure(figsize=(12,6))
plt.title(schema.Gender)
```

```
plt.pie(gender_counts, labels=gender_counts.index, autopct='%1.1f%%', startangle=180);
```

Which of the following describe you, if any? Please check all that apply. If you prefer not to answer, you may leave this question blank.



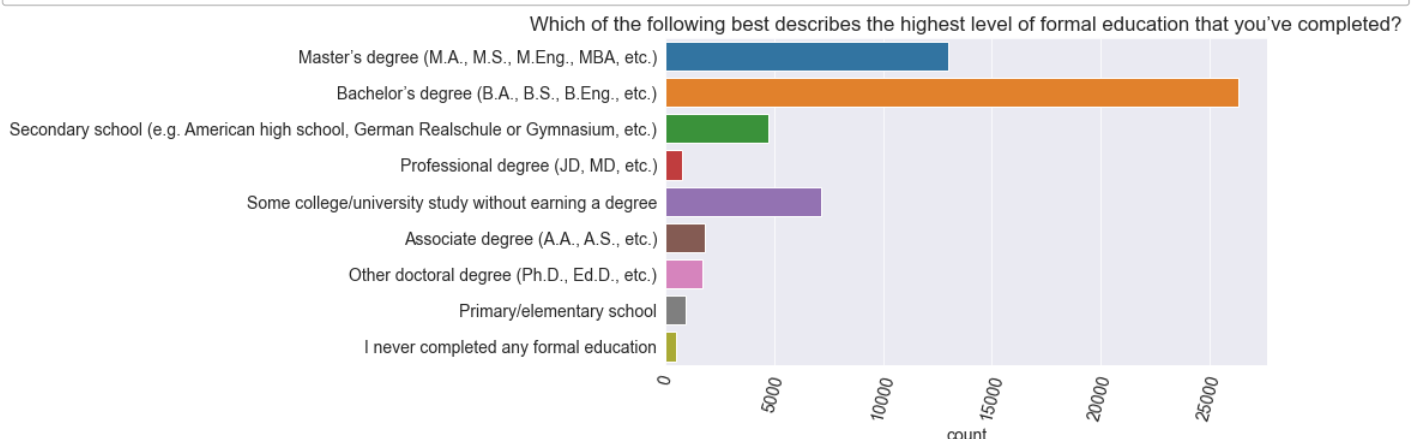
Only about 8% of survey respondents who have answered the question identify as women or non-binary. This number is lower than the overall percentage of women & non-binary genders in the programming community - which is estimated to be around 12%.

Exercise: It would be interesting to compare the survey responses & preferences across genders. Repeat this analysis with these breakdowns. How do the relative education levels differ across genders? How do the salaries vary? You may find this analysis on the [Gender Divide in Data Science](#) useful.

Education Level

Formal education in computer science is often considered an essential requirement for becoming a programmer. However, there are many free resources & tutorials available online to learn programming. Let's compare the education levels of respondents to gain some insight into this. We'll use a horizontal bar plot here.

```
sns.countplot(y=survey_df.EdLevel)
plt.xticks(rotation=75);
plt.title(schema['EdLevel'])
plt.ylabel(None);
```



It appears that well over half of the respondents hold a bachelor's or master's degree, so most programmers seem to have some college education. However, it's not clear from this graph alone if they hold a degree in computer science.

Exercises: The graph currently shows the number of respondents for each option. Can you modify it to show the percentage instead? Further, try comparing the percentages for each degree for men vs. women.

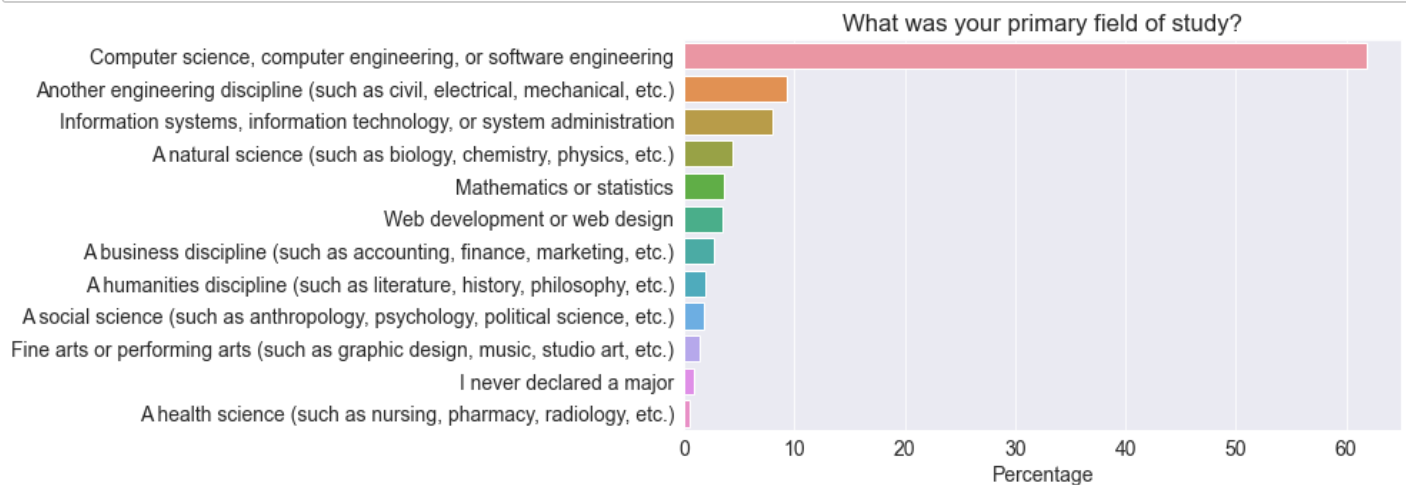
Let's also plot undergraduate majors, but this time we'll convert the numbers into percentages and sort the values to make it easier to visualize the order.

```
schema.UndergradMajor
```

```
'What was your primary field of study?'
```

```
undergrad_pct = survey_df.UndergradMajor.value_counts() * 100 / survey_df.UndergradMajor
sns.barplot(x=undergrad_pct, y=undergrad_pct.index)

plt.title(schema.UndergradMajor)
plt.ylabel(None);
plt.xlabel('Percentage');
```



It turns out that 40% of programmers holding a college degree have a field of study other than computer science - which is very encouraging. It seems to suggest that while a college education is helpful in general, you do not need to pursue a major in computer science to become a successful programmer.

Exercises: Analyze the `NEWEdImpt` column for respondents who hold some college degree vs. those who don't. Do you notice any difference in opinion?


Employment

Freelancing or contract work is a common choice among programmers, so it would be interesting to compare the breakdown between full-time, part-time, and freelance work. Let's visualize the data from the `Employment` column.

```
schema.Employment
```

```
'Which of the following best describes your current employment status?'
```

```
(survey_df.Employment.value_counts(normalize=True, ascending=True)*100).plot(kind='barh')
plt.title(schema.Employment)
plt.xlabel('Percentage');
```



It appears that close to 10% of respondents are employed part time or as freelancers.

Exercise: Add a new column `EmploymentType` containing the values `Enthusiast` (student or not employed but looking for work), `Professional` (employed full-time, part-time or freelancing), and `Other` (not employed or retired). For each of the graphs that follow, show a comparison between `Enthusiast` and `Professional`.

The `DevType` field contains information about the roles held by respondents. Since the question allows multiple answers, the column contains lists of values separated by a semi-colon `;`, making it a bit harder to analyze directly.

```
schema.DevType
```

'Which of the following describe you? Please select all that apply.'

```
survey_df.DevType.value_counts()
```

```
Developer, full-stack
4396
Developer, back-end
3056
Developer, back-end;Developer, front-end;Developer, full-stack
2214
Developer, back-end;Developer, full-stack
1465
Developer, front-end
1390
...
Developer, back-end;Developer, embedded applications or devices;Developer, front-
end;Developer, full-stack;Developer, game or graphics;Engineering manager;Senior
executive/VP
1
Data or business analyst;Data scientist or machine learning specialist;Developer,
mobile;Developer, QA or test;Product manager
1
Developer, full-stack;Developer, QA or test;Product manager
1
```


Name: DevType, Length: 8213, dtype: int64

```
def split_multicolumn(col_series):
    result_df = col_series.to_frame()
    options = []
    # Iterate over the column
    for idx, value in col_series[col_series.notnull()].iteritems():
        # Break each value into list of options
        for option in value.split(';'):
            # Add the option as a column to result
            if not option in result_df.columns:
                options.append(option)
                result_df[option] = False
            # Mark the value in the option column as True
            result_df.at[idx, option] = True
    return result_df[options]
```

```
dev_type_df = split_multicolumn(survey_df.DevType)
```

```
dev_type_df
```

[illegible]

64306 rows × 23 columns

The `dev_type_df` has one column for each option that can be selected as a response. If a respondent has chosen an option, the corresponding column's value is `True`. Otherwise, it is `False`.

We can now use the column-wise totals to identify the most common roles.

```
dev_type_totals = dev_type_df.sum().sort_values(ascending=False)
dev_type_totals
```

Developer, back-end	26996
Developer, full-stack	26915
Developer, front-end	18128
Developer, desktop or enterprise applications	11687
Developer, mobile	9406
DevOps specialist	5915
Database administrator	5658
Designer	5262
System administrator	5185
Developer, embedded applications or devices	4701
Data or business analyst	3970
Data scientist or machine learning specialist	3939
Developer, QA or test	3893
Engineer, data	3700
Academic researcher	3502
Educator	2895
Developer, game or graphics	2751
Engineering manager	2699
Product manager	2471
Scientist	2060
Engineer, site reliability	1921
Senior executive/VP	1292
Marketing or sales professional	625

dtype: int64

As one might expect, the most common roles include "Developer" in the name.

Exercises:

- Can you figure out what percentage of respondents work in roles related to data science?
- Which positions have the highest percentage of women?

We've only explored a handful of columns from the 20 columns that we selected. Explore and visualize the remaining columns using the empty cells below.

Let's save and upload our work before continuing.

```
import jovian
```

```
jovian.commit()
```

```
[jovian] Attempting to save notebook..
```

```
[jovian] Updating notebook "aakashns/python-eda-stackoverflow-survey" on  
https://jovian.ai/
```

```
[jovian] Uploading notebook..
```

```
[jovian] Capturing environment..
```

```
[jovian] Committed successfully! https://jovian.ai/aakashns/python-eda-stackoverflow-survey
```

```
'https://jovian.ai/aakashns/python-eda-stackoverflow-survey'
```

Asking and Answering Questions

We've already gained several insights about the respondents and the programming community by exploring individual columns of the dataset. Let's ask some specific questions and try to answer them using data frame operations and visualizations.

Q: What are the most popular programming languages in 2020?

To answer, this we can use the `LanguageWorkedWith` column. Similar to `DevType`, respondents were allowed to choose multiple options here.

```
survey_df.LanguageWorkedWith
```

```
0          C#;HTML/CSS;JavaScript
1          JavaScript;Swift
2      Objective-C;Python;Swift
3                          NaN
4          HTML/CSS;Ruby;SQL
...
64456                          NaN
64457  Assembly;Bash/Shell/PowerShell;C;C#;C++;Dart;G...
64458                          NaN
64459                          HTML/CSS
64460      C#;HTML/CSS;Java;JavaScript;SQL
Name: LanguageWorkedWith, Length: 64306, dtype: object
```

First, we'll split this column into a data frame containing a column of each language listed in the options.

```
languages_worked_df = split_multicolumn(survey_df.LanguageWorkedWith)
```

```
languages_worked_df
```

	C#	HTML/CSS	JavaScript	Swift	Objective-C	Python	Ruby	SQL	Java	PHP	...	VBA	Perl	Scala
0	True	True	True	False	False	False	False	False	False	False	...	False	False	False
1	False	False	True	True	False	False	False	False	False	False	...	False	False	False
2	False	False	False	True	True	True	False	False	False	False	...	False	False	False
3	False	False	False	False	False	False	False	False	False	False	...	False	False	False
4	False	True	False	False	False	False	True	True	False	False	...	False	False	False
...
64456	False	False	False	False	False	False	False	False	False	False	...	False	False	False
64457	True	True	True	True	True	True	True	True	True	True	...	True	True	True
64458	False	False	False	False	False	False	False	False	False	False	...	False	False	False
64459	False	True	False	False	False	False	False	False	False	False	...	False	False	False
64460	True	True	True	False	False	False	False	True	True	False	...	False	False	False

64306 rows × 25 columns

It appears that a total of 25 languages were included among the options. Let's aggregate these to identify the percentage of respondents who selected each language.

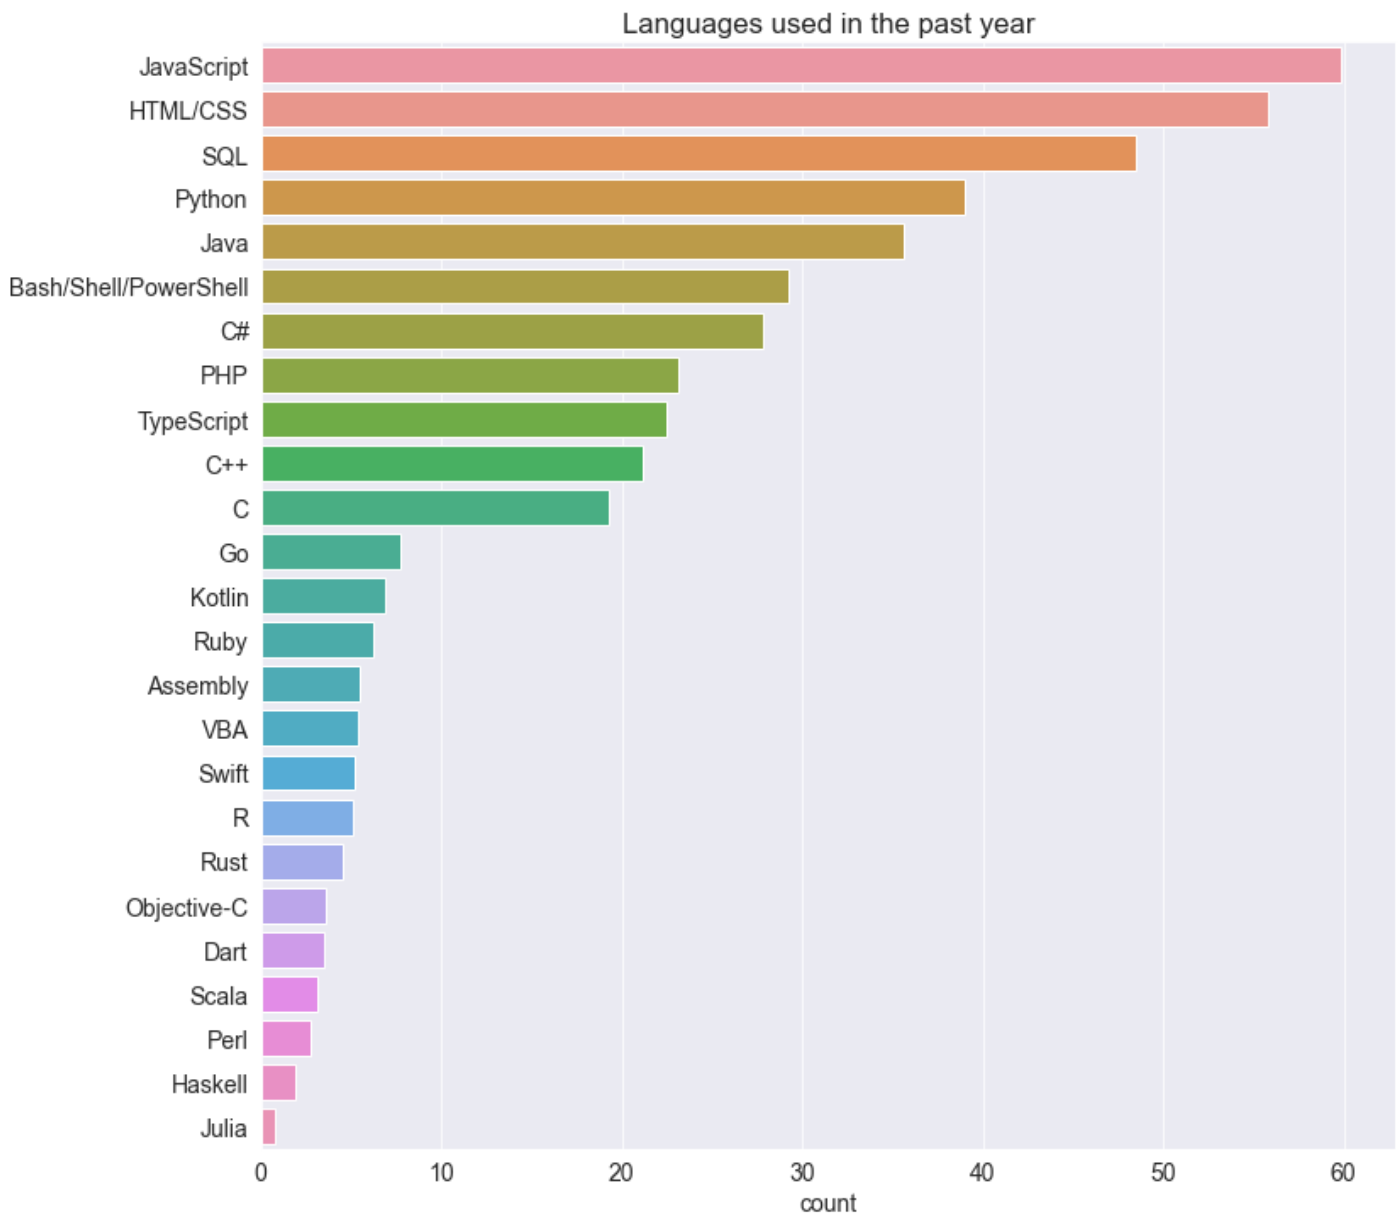
```
languages_worked_percentages = languages_worked_df.mean().sort_values(ascending=False)
languages_worked_percentages
```

```
JavaScript          59.893323
HTML/CSS           55.801947
SQL                 48.444935
Python              39.001026
Java                35.618760
Bash/Shell/PowerShell 29.239884
C#                  27.803004
PHP                 23.130035
TypeScript          22.461357
C++                 21.114670
C                   19.236152
Go                  7.758219
Kotlin              6.887382
Ruby                6.229590
Assembly            5.447392
VBA                 5.394520
Swift               5.226573
R                   5.064846
Rust                4.498803
Objective-C         3.603085
Dart                3.517557
Scala               3.150561
Perl                2.757130
```

```
Haskell          1.861413
Julia             0.782198
dtype: float64
```

We can plot this information using a horizontal bar chart.

```
plt.figure(figsize=(12, 12))
sns.barplot(x=languages_worked_percentages, y=languages_worked_percentages.index)
plt.title("Languages used in the past year");
plt.xlabel('count');
```



Perhaps unsurprisingly, Javascript & HTML/CSS comes out at the top as web development is one of today's most sought skills. It also happens to be one of the easiest to get started. SQL is necessary for working with relational databases, so it's no surprise that most programmers work with SQL regularly. Python seems to be the popular choice for other forms of development, beating out Java, which was the industry standard for server & application development for over two decades.

Exercises:

- What are the most common languages used by students? How does the list compare with the most common languages used by professional developers?

- What are the most common languages among respondents who do not describe themselves as "Developer, front-end"?
- What are the most common languages among respondents who work in fields related to data science?
- What are the most common languages used by developers older than 35 years of age?
- What are the most common languages used by developers in your home country?

Q: Which languages are the most people interested to learn over the next year?

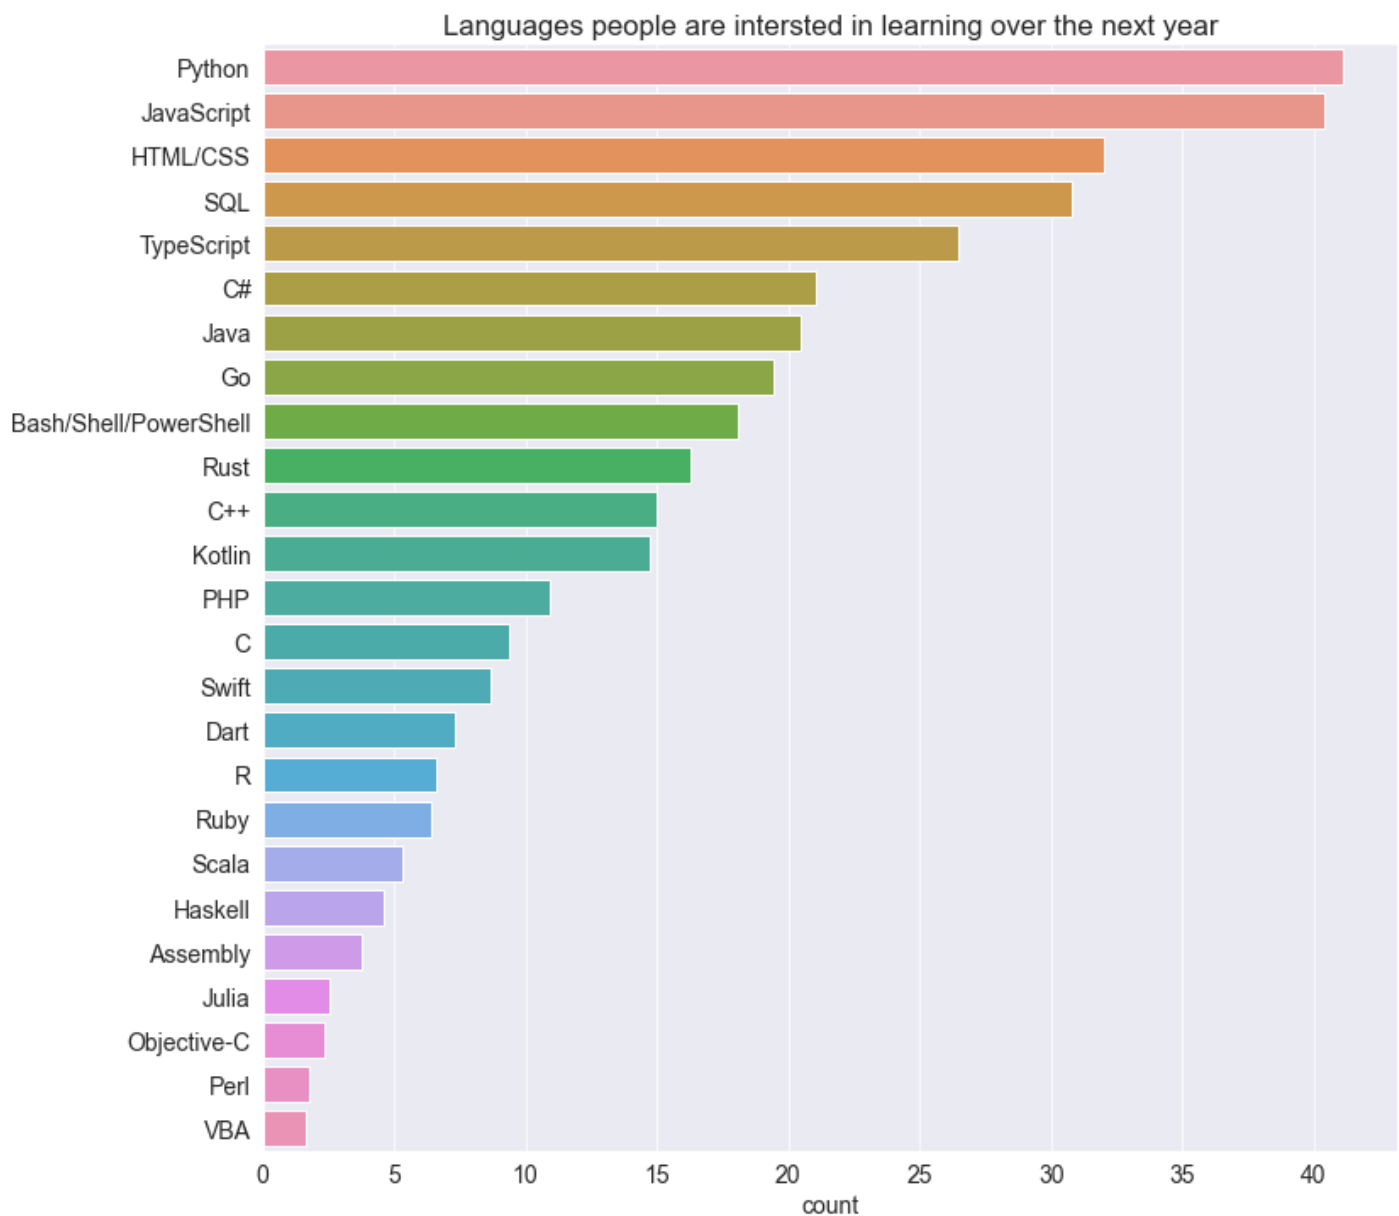
For this, we can use the `LanguageDesireNextYear` column, with similar processing as the previous one.

```
languages_interested_df = split_multicolumn(survey_df.LanguageDesireNextYear)
languages_interested_percentages = languages_interested_df.mean().sort_values(ascending=False)
languages_interested_percentages
```

Python	41.143906
JavaScript	40.425466
HTML/CSS	32.028116
SQL	30.799614
TypeScript	26.451653
C#	21.058688
Java	20.464653
Go	19.432090
Bash/Shell/PowerShell	18.057413
Rust	16.270643
C++	15.014151
Kotlin	14.760676
PHP	10.947657
C	9.359935
Swift	8.692812
Dart	7.308805
R	6.571704
Ruby	6.425528
Scala	5.326097
Haskell	4.593662
Assembly	3.766367
Julia	2.540976
Objective-C	2.338818
Perl	1.761888
VBA	1.611047

dtype: float64

```
plt.figure(figsize=(12, 12))
sns.barplot(x=languages_interested_percentages, y=languages_interested_percentages.index)
plt.title("Languages people are intersted in learning over the next year");
plt.xlabel('count');
```



Once again, it's not surprising that Python is the language most people are interested in learning - since it is an easy-to-learn general-purpose programming language well suited for a variety of domains: application development, numerical computing, data analysis, machine learning, big data, cloud automation, web scraping, scripting, etc. We're using Python for this very analysis, so we're in good company!

Exercises: Repeat the exercises from the previous question, replacing "most common languages" with "languages people are interested in learning/using."

Q: Which are the most loved languages, i.e., a high percentage of people who have used the language want to continue learning & using it over the next year?


While this question may seem tricky at first, it's straightforward to solve using Pandas array operations. Here's what we can do:

- Create a new data frame `languages_loved_df` that contains a `True` value for a language only if the corresponding values in `languages_worked_df` and `languages_interested_df` are both `True`
- Take the column-wise sum of `languages_loved_df` and divide it by the column-wise sum of `languages_worked_df` to get the percentage of respondents who "love" the language
- Sort the results in decreasing order and plot a horizontal bar graph

```
languages_loved_df = languages_worked_df & languages_interested_df
```

```
languages_loved_percentages = (languages_loved_df.sum() * 100 / languages_worked_df.sum()
```

```
plt.figure(figsize=(12, 12))
sns.barplot(x=languages_loved_percentages, y=languages_loved_percentages.index)
plt.title("Most loved languages");
plt.xlabel('count');
```



[Rust](#) has been StackOverflow's most-loved language for [four years in a row](#). The second most-loved language is TypeScript, a popular alternative to JavaScript for web development.

Python features at number 3, despite already being one of the most widely-used languages in the world. Python has a solid foundation, is easy to learn & use, has a large ecosystem of domain-specific libraries, and a massive worldwide community.

Exercises: What are the most dreaded languages, i.e., languages which people have used in the past year but do not want to learn/use over the next year. Hint: `~languages_interested_df`.

Q: In which countries do developers work the highest number of hours per week?
Consider countries with more than 250 responses only.

To answer this question, we'll need to use the `groupby` data frame method to aggregate the rows for each country. We'll also need to filter the results to only include the countries with more than 250 respondents.

```
countries_df = survey_df.groupby('Country')[['WorkWeekHrs']].mean().sort_values('WorkWeekHrs')
```

```
high_response_countries_df = countries_df.loc[survey_df.Country.value_counts() > 250].h
```

```
high_response_countries_df
```

WorkWeekHrs	
Country	
Iran	44.337748
Israel	43.915094
China	42.150000
United States	41.802982
Greece	41.402724
Viet Nam	41.391667
South Africa	41.023460
Turkey	40.982143
Sri Lanka	40.612245
New Zealand	40.457551
Belgium	40.444444
Canada	40.208837
Hungary	40.194340
Bangladesh	40.097458
India	40.090603

The Asian countries like Iran, China, and Israel have the highest working hours, followed by the United States. However, there isn't too much variation overall, and the average working hours seem to be around 40 hours per week.

Exercises:

- How do the average work hours compare across continents? You may find this list of [countries in each continent](#) useful.
- Which role has the highest average number of hours worked per week? Which one has the lowest?
- How do the hours worked compare between freelancers and developers working full-time?

Q: How important is it to start young to build a career in programming?

Let's create a scatter plot of `Age` vs. `YearsCodePro` (i.e., years of coding experience) to answer this question.

```
schema.YearsCodePro
```

'NOT including education, how many years have you coded professionally (as a part of your work)?'

```
sns.scatterplot(x='Age', y='YearsCodePro', hue='Hobbyist', data=survey_df)
plt.xlabel("Age")
plt.ylabel("Years of professional coding experience");
```

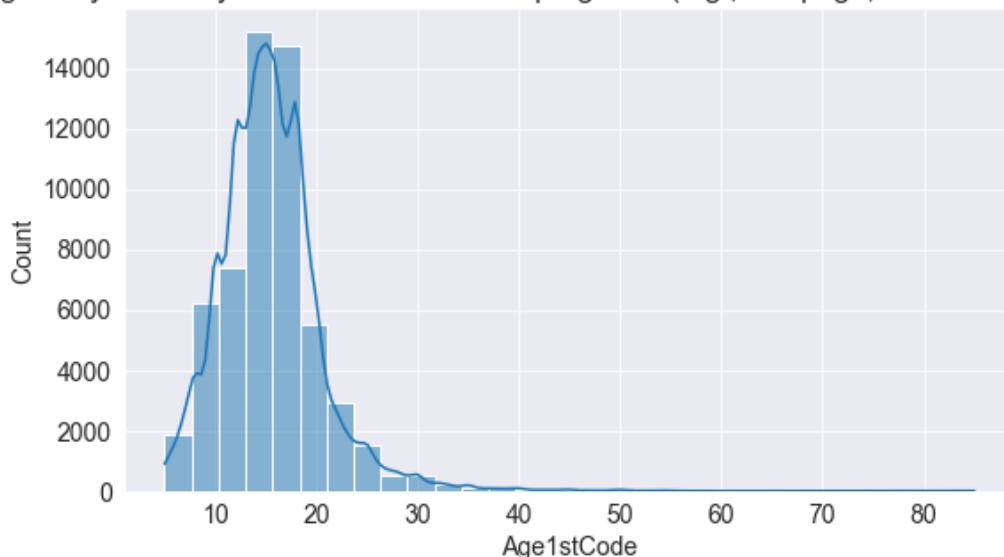


You can see points all over the graph, which indicates that you can **start programming professionally at any age**. Many people who have been coding for several decades professionally also seem to enjoy it as a hobby.

We can also view the distribution of the `Age1stCode` column to see when the respondents tried programming for the first time.

```
plt.title(schema.Age1stCode)
sns.histplot(x=survey_df.Age1stCode, bins=30, kde=True);
```

At what age did you write your first line of code or program? (e.g., webpage, Hello World, Scratch project)



As you might expect, most people seem to have had some exposure to programming before the age of 40. However, but there are people of all ages and walks of life learning to code.

Exercises:

- How does programming experience change opinions & preferences? Repeat the entire analysis while comparing the responses of people who have more than ten years of professional programming experience vs. those who don't. Do you see any interesting trends?
- Compare the years of professional coding experience across different genders.

Hopefully, you are already thinking of many more questions you'd like to answer using this data. Use the empty cells below to ask and answer more questions.

Let's save and commit our work before continuing

```
import jovian
```

```
jovian.commit()
```

```
[jovian] Attempting to save notebook..
```

```
[jovian] Updating notebook "aakashns/python-eda-stackoverflow-survey" on  
https://jovian.ai/
```

```
[jovian] Uploading notebook..
```

```
[jovian] Capturing environment..
```

```
[jovian] Committed successfully! https://jovian.ai/aakashns/python-eda-stackoverflow-survey
```

```
'https://jovian.ai/aakashns/python-eda-stackoverflow-survey'
```

Inferences and Conclusions

We've drawn many inferences from the survey. Here's a summary of a few of them:

- Based on the survey respondents' demographics, we can infer that the survey is somewhat representative of the overall programming community. However, it has fewer responses from programmers in non-English-speaking countries and women & non-binary genders.
- The programming community is not as diverse as it can be. Although things are improving, we should make more efforts to support & encourage underrepresented communities, whether in terms of age, country, race,

gender, or otherwise.

- Although most programmers hold a college degree, a reasonably large percentage did not have computer science as their college major. Hence, a computer science degree isn't compulsory for learning to code or building a career in programming.
- A significant percentage of programmers either work part-time or as freelancers, which can be a great way to break into the field, especially when you're just getting started.
- Javascript & HTML/CSS are the most used programming languages in 2020, closely followed by SQL & Python.
- Python is the language most people are interested in learning - since it is an easy-to-learn general-purpose programming language well suited for various domains.
- Rust and TypeScript are the most "loved" languages in 2020, both of which have small but fast-growing communities. Python is a close third, despite already being a widely used language.
- Programmers worldwide seem to be working for around 40 hours a week on average, with slight variations by country.
- You can learn and start programming professionally at any age. You're likely to have a long and fulfilling career if you also enjoy programming as a hobby.

Exercises

There's a wealth of information to be discovered using the survey, and we've barely scratched the surface. Here are some ideas for further exploration:

- Repeat the analysis for different age groups & genders, and compare the results
- Pick a different set of columns (we chose 20 out of 65) to analyze other facets of the data
- Prepare an analysis focusing on diversity - and identify areas where underrepresented communities are at par with the majority (e.g., education) and where they aren't (e.g., salaries)
- Compare the results of this year's survey with the previous years and identify interesting trends

References and Future Work

Check out the following resources to learn more about the dataset and tools used in this notebook:

- Stack Overflow Developer Survey: <https://insights.stackoverflow.com/survey>
- Pandas user guide: https://pandas.pydata.org/docs/user_guide/index.html
- Matplotlib user guide: <https://matplotlib.org/3.3.1/users/index.html>
- Seaborn user guide & tutorial: <https://seaborn.pydata.org/tutorial.html>
- opendatasets Python library: <https://github.com/JovianML/opendatasets>

As a next step, you can try out a project on another dataset of your choice:

<https://jovian.ml/aakashns/zerotopandas-course-project-starter> .

```
import jovian
```

```
jovian.commit()
```

[jovian] Attempting to save notebook..